

Aravindh MS — SRE & Production Operations

A resume-grounded preparation companion: ten years of production support, automation and reliability work across FinTech, Healthcare IT and Entertainment, organised into stories, deep-dives and talking points.

SRE

Application Support

Technical Lead

Production Operations



The operating floor where ten years of incident response, automation and release work come together.

At a glance

This opening page sets the scene for the full cheatsheet: Aravindh's 10+ years spanning Cognizant, PagoNxt, Citi and TAO, the four target roles this guide prepares for, and how the remaining nine pages build from positioning through STAR stories, technical deep-dives and role-specific question banks — all anchored to the five quantified wins on the resume.

Inside this guide

02

Elevator Pitch & Positioning

03

STAR Stories

04

Behavioral Q&A

05

Technical Deep-Dives

06

Role-Specific Question Banks

Positioning & Career Narrative

Aravindh MS · SRE / Application Support / Technical Lead / Production Operations · Nature Vision · Sep 20, 2026



30-SECOND PITCH

"I'm a production operations engineer with 10+ years keeping payments, healthcare and entertainment platforms up — from 24×7 monitoring at Cognizant to Blue-Green releases at PagoNxt and DR planning at Citi. I automate first, escalate second."


60-SECOND PITCH

Start with the 30-second frame, then anchor one proof point: at PagoNxt I rebuilt monitoring and alerting and cut incident response time by 40%. At Cognizant, Shell/Python automation with crontab removed 30% of manual support effort. Close with the role you want next.

2-MINUTE PITCH — SKELETON

- **Who I am:** L2/L3 support and SRE-leaning engineer across FinTech, Healthcare IT and Entertainment.
- **How I work:** investigate at the API/DB/log layer (Control-M, Informix/SQL), then fix the system, not the ticket.
- **Proof:** -40% incident response (PagoNxt), -40% downtime incidents (Cognizant), -30% manual support (Cognizant).
- **Leadership:** mentoring, SOP building and release coordination — the Technical Lead story.
- **Why this role:** tie the target team's stack to the tooling you already run daily.

The Career Arc — One Sentence Per Stop



Narrative thread: every move added a layer — reliability tooling, then resilience discipline, then people leadership.

Summary

Lead every interview answer with the same through-line: a decade of keeping critical systems running, hardened by automation and measurable outcomes. Pick the pitch length that matches the moment — 30 seconds for openers, 60 for recruiter screens, the 2-minute skeleton for hiring managers — and always land on one quantified win.

The 40% Response-Time Rescue

- 01 Story breakdown
- 02 Result metric
- 03 Interview follow-ups
- 04 Why it lands

Monitoring & alerting overhaul for payments operations

- S
- Payment platform incidents at PagoNxt were surfacing late — alerts reached on-call engineers after customer impact had already spread across services.
- T
- Rebuild the monitoring and alerting layer so production issues were detected and routed before they escalated, without flooding the team with noise.
- A
- Tuned alert thresholds and severity routing, wired service health checks into the on-call workflow, and coordinated release windows with the Blue-Green deployment process so changes and alerts stopped colliding.
- R
- Incident response time dropped sharply and on-call handovers became predictable, with cleaner escalation paths between L1, L2 and release teams.

- 40% incident response time — the headline metric to quote in any SRE or Production Operations interview.

Follow-up questions to expect

- How did you decide which alerts were actionable vs. noise?
- What did the escalation path look like before and after?
- How did Blue-Green releases interact with alert routing?




On-call production operations — the setting behind the response-time rescue.

Why it lands: one metric, one system, one owner — lead with the -40% number, then walk the interviewer through Situation → Task → Action → Result in under two minutes.

The 40% Response-Time Rescue

A resume-grounded interview story: how a monitoring and alerting redesign at PagoNxt cut incident response time by 40%. Reach for this story when asked about observability, ownership under pressure, or production firefighting.



S • SITUATION

Payment flows at PagoNxt were failing intermittently; alerts reached on-call engineers late, and every triage started from scratch.

T • TASK


Own the monitoring and alerting layer so incidents are detected and routed faster — without adding headcount.

A • ACTION

Reworked alert thresholds and routing, tied alerts to payment-path components, and built runbook links so responders skipped discovery.

R • RESULT

Incident response time dropped **-40%**, with cleaner L1 → L2 handoffs and fewer repeat escalations.



Interview takeaway

The -40% response-time gain came not from new tooling spend but from re-mapping alerts to payment-path components and embedding runbooks in the alert itself — a pattern that transfers directly to Datadog, Splunk, or ELK environments.

Blue-Green Confidence

How **Blue-Green deployments and coordinated release management** at PagoNxt turned risky payment-platform releases into low-drama, reversible events — the story I tell when interviewers ask about deployment safety.

The story, in four beats

- SIT

Payment service releases carried customer-visible risk; a failed deployment meant rollback under pressure during live transaction windows.
- TASK

Own release coordination so every deployment had a tested fallback path and a clear go/no-go owner.
- ACT


Introduced Blue-Green environments — parallel stacks with instant traffic switch — plus pre-release checklists, smoke tests, and a timed cutover runbook agreed with stakeholders.
- RESULT

Releases became reversible in minutes, not hours; rollback anxiety disappeared from release calls and coordination overhead dropped.



Release night, done calmly: parallel environments standing by before cutover.

Cutover mechanics



What made it land

- Minutes-level rollback
- Zero-drama release calls
- Smoke tests before cutover
- Single go/no-go owner

Follow-up questions I prepare for

- How do database migrations behave when two versions run side by side?
- What triggers an automatic vs. human rollback decision?
- How does this extend to canary or progressive delivery?

SUMMARY

Blue-Green deployments plus disciplined release coordination made every PagoNxt deployment reversible in minutes — the difference between a release event and a release routine.

DR Test Day Lead at Citi

At Citi, Aravindh led disaster-recovery test days end to end — coordinating with data-centre operations to fail workloads over, validate recovery objectives, and bring services back without surprises.



DR test day: rehearsed failover under real production conditions.

Situation

Citi required periodic BCP/DR exercises proving that critical applications could survive a data-centre loss. Past drills were treated as paperwork — runbooks were stale and recovery steps untested.

Task


Own the DR test plan for supported applications: define failover scope, agree RTO/RPO targets with stakeholders, and coordinate execution with DC operations within a fixed change window.

Action

Refreshed runbooks step by step, sequenced DNS and traffic cutover, rehearsed database recovery and application bring-up in order, and staffed a live bridge to log every checkpoint and rollback trigger.

Result

Failover completed inside the agreed recovery window with verified data integrity, and the exercised runbooks became the standing reference for subsequent BCP/DR cycles — no improvised steps on the day.



Likely follow-up questions

- How did you set a realistic RTO, and what did you do when a step threatened to breach it?
- What was your rollback trigger, and who had the authority to call it?
- How did you keep runbooks from going stale between drills?

Five Numbers That Carry Every Answer

Each quantified win below is a ready-made proof point. Anchor a story with one, then let the interviewer pull the thread.



Metric-first answers start on the operations floor — dashboards, alert tiers, and a single headline number.

Talking points by metric

Metric	Where it happened	How to weave it in
-40%	Incident response time PagoNxt • monitoring & alerting redesign	Use when asked about observability or alert fatigue — lead with the number, then explain the alert-tiering logic behind it.
-40%	Downtime incidents Cognizant • 24×7 job monitoring	Best fit for reliability and proactive-monitoring questions; pair with the overnight batch-watch routine.
-30%	Manual support effort Cognizant • Shell/Python automation with crontab	Deploy in automation questions — describe one script that replaced a manual runbook step.
+30%	Operational performance Cross-team process improvements	Frame as an outcome of tuning workflows, not a single tool — shows systems thinking.
+20%	Team productivity Mentoring, SOPs & knowledge sharing at Citi	Use for leadership questions; tie to SOP adoption and onboarding speed for new engineers.

Summary: one number per answer

Pick a single metric per interview answer and state it early — “cut incident response time by 40%” — then spend the rest of the story on how. Never stack two metrics in one sentence; it dilutes both and invites scepticism.

Memory Hooks & the 10-Minute Pre-Interview Run

Everything on this page is built for the final minutes before the call: one-line hooks that unlock each STAR story, a timed checklist, and the do/don't guardrails.



The last quiet ten minutes — hooks first, panic never.

One-Line Memory Hooks

PagoNxt
"Monitoring rebuilt, response cut 40%." Alert redesign → faster triage → -40% incident response time.

Cognizant · Downtime
"24×7 eyes on jobs, downtime down 40%." Proactive batch monitoring cut repeat incidents.

Cognizant · Automation
"Shell + Python + crontab = -30% manual work." Routine support automated, engineers freed for real problems.

PagoNxt · Releases
"Blue-Green confidence." Zero-drama deployments with instant rollback and release coordination.

Citi · DR
"DR Day Lead." Ran DR testing and traffic management with DC ops; BCP/DR planning owner.

Citi · Mentoring
"SOPs + mentoring = +20% team productivity." Documented runbooks, coached juniors, stabilized KNIME workflows.

The 10-Minute Pre-Interview Checklist

- ✓ **Min 0–2:** Say the 30-second pitch out loud once — Cognizant → PagoNxt → Citi → TAO arc.
- ✓ **Min 2–4:** Recite all six hooks above from memory, no peeking.
- ✓ **Min 4–6:** Re-read the two stories most relevant to this role (SRE → monitoring; Prod Ops → DR Day).
- ✓ **Min 6–7:** Refresh the five numbers: -40%, -40%, -30%, +30%, +20%.
- ✓ **Min 7–8:** Skim the role-specific question bank for this interview round.
- ✓ **Min 8–10:** Water, notepad, two questions to ask — then breathe.

Do & Don't in the Room

- Do**
- Anchor every answer to a real system you ran — Control-M, Rundeck, Terraform, Jenkins, Datadog.
 - Lead with the result, then unpack Situation → Task → Action.
 - Quantify: "cut response time 40%" beats "improved monitoring".
 - Name the ITIL practice you followed — incident, change, problem, request.

- Don't**
- Recite the resume line by line — tell the story behind it.
 - Claim tools you can't defend one follow-up deep.
 - Blame past teams when discussing failures; own the fix.
 - Let a silence fill — "Let me give you a concrete example" is your bridge.

Summary: Revision is retrieval, not re-reading. Six hooks, five numbers, two role-matched stories, and a calm first minute — that is the whole pre-interview playbook.

Speak the Language, Then Flip the Interview

A rapid-fire vocabulary refresher for SRE, ITIL and DevOps terms — followed by smart questions Aravindh can ask in each interview round to signal seniority.

Key Terms, Crisp Definitions

SLO / SLA / SLI

Service Level Objective (internal target), Agreement (customer contract), Indicator (the measured signal). SRE negotiates SLOs from SLIs.

Error Budget

1 – SLO: the tolerated unreliability. Exhausted budget freezes feature releases in favour of reliability work.

RT0 / RP0

Recovery Time Objective — how fast service must return; Recovery Point Objective — maximum tolerable data loss window. Both anchor DR testing.

Blue-Green Deployment

Two identical environments; traffic switches atomically from blue to green, enabling instant rollback — used for release coordination at PagoNxt.

MTTR / MTTD

Mean Time to Resolve / Detect. Monitoring and alerting redesigns directly shrink both — the basis of the -40% response-time win.

ITIL Incident vs Problem

Incident = restore service fast; Problem = eliminate root cause so incidents stop recurring.

IaC (Terraform)

Infrastructure declared as versioned code — reviewable, repeatable, drift-detectable across environments.

Control-M / Rundeck

Enterprise job schedulers; batch dependency graphs and runbook automation for 24×7 operations.

Questions to Ask, by Round

Technical round

- "What do your SLOs and error budgets look like today?"
- "How much of your infrastructure is Terraform-managed vs manual?"
- "What does your worst recurring alert look like, and who owns it?"

Hiring manager round

- "What separates a good on-call engineer from a great one here?"
- "How does the team balance feature work against reliability debt?"
- "Where did the last postmortem change an actual process?"

HR / culture round

- "How is on-call compensation and rotation fairness handled?"
- "What does success look like at 90 days for this role?"
- "How are mentoring and SOP ownership recognised?"



Preparation is a two-way conversation: vocabulary in hand, questions ready.

8

core terms rehearsed until definitions come out in one clean sentence

9

questions to ask, tiered across technical, manager and HR rounds


1

closing move per round that ties a question to the interviewer's own stack

Summary: this page closes the cheatsheet's communication layer — a memorisable glossary of the terms interviewers expect, plus round-specific questions that demonstrate fluency rather than recitation. Used with the pitch and STAR pages, it converts vocabulary into a two-way conversation that lands as seniority.

One Guide, Ready When You Are

Ten pages in, the preparation loop is complete: pitch, stories, deep-dives, question banks, metrics, revision hooks, and questions to ask. Everything traces back to real work at Cognizant, PagoNxt, Citi and TAO — nothing to memorise that isn't already yours.



4-step loop · repeat per interview round

The readiness loop: revise, answer, ask, learn — each pass sharpens the next round.

Sections to revisit

- **Elevator pitches** — 30s / 60s / 2-min versions.
- **STAR library** — 6–8 resume-grounded stories.
- **Technical deep-dives** — tooling, DR, deployments.
- **Role question banks** — SRE, App Support, Lead, Prod Ops.
- **Rapid revision** — ten-minute pre-call scan.



End of shift, start of the next opportunity.

Summary — the takeaway

Preparation isn't about knowing more — it's about retrieving faster. This cheatsheet turns ten years of production operations, SRE practice and technical leadership into answers you can reach in seconds. Run the readiness loop before every round: open with a number, match the STAR story to the question, stay resume-grounded, and close by asking one question back. Let the metrics do the persuading.